

Predictive performance across scales: raw node set versus distribution-matched core

